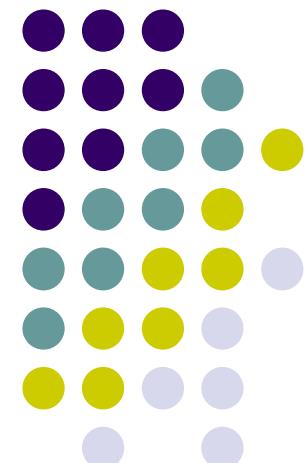
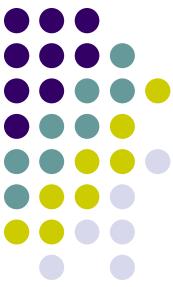


# *Mastering Data Structures and Algorithms: A Practical Approach*

## Greedy Algorithms

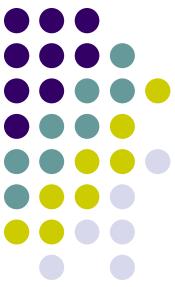


By Dr. Juan C. Gomez  
Fall 2018



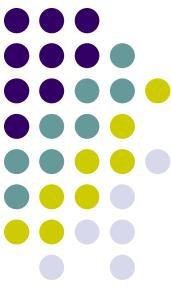
# Overview

- Greedy Algorithms: what are they?
- Properties of Solution
- Concept
- Classical Problems
  - Minimum Spanning Tree
    - Prim's Algorithm
    - Kruskal's Algorithm
  - Shortest path
    - Dijkstra's Algorithm
  - Huffman Trees
    - Huffman Codes



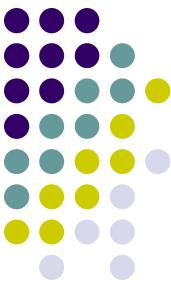
# Overview

- Classical Problems
  - Activity Selection Problem



# Greedy Algorithms

- Illustrate by using the “Change Making Problem”:
  - Given you have coins that are of the following denomination:
    - 25c, 10c, 5c, and 1c
    - Provide change for a given amount: i.e. 48c
  - We use the highest denomination first (we want to provide a solution with the least number of coins:
    - Give 25c, we can't give 2 25c as that will violate constraints of our required solution.
    - Now we add 2 x 10c ( $25c + 2 \times 10c < 48c$ )
    - We can't add any 5c as that will violate the constraints of our required solution.
    - Now we add 3 x 1c ( $25c + 2 \times 10c + 3 \times 1c = 48c$ )
    - Characteristics of the solution: optimal, less number of coins!



# Greedy Algorithms

- Always make the choice that seems best at the moment.
  - Local optimal choice leads to Global Optimal
- Given an optimization problem
  - It is usually easy to find a solution analyzing small-size problem.
  - It is not as easy to prove that the solution given by Greedy approach is optimal.
    - Usually, induction may be used to prove solution is optimal.



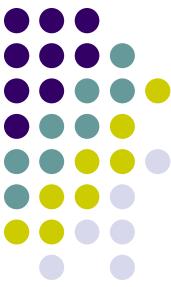
# Properties of solutions

- At every step ensure the following:
  - *feasible*, i.e., it has to satisfy the problem's constraints
  - *locally optimal*, i.e., it has to be the best local choice among all feasible choices available on that step
  - *irrevocable*, i.e., once made, it cannot be changed on subsequent steps of the algorithm



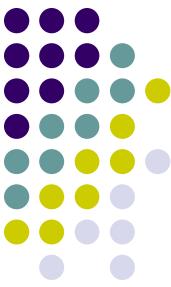
# Concept

- At every step choose a local optimal solution.
- Hope that this combination of locally optimal solutions will lead to a Global Optimum solution.
- Works for some problems, but not for others.



# Minimum Spanning Tree

The following problem arises naturally in many practical situations: given  $n$  points, connect them in the cheapest possible way so that there will be a path between every pair of points. It has direct applications to the design of all kinds of networks—including communication, computer, transportation, and electrical—by providing the cheapest way to achieve connectivity. It identifies clusters of points in data sets. It has been used for classification purposes in archeology, biology, sociology, and other sciences. It is also helpful for constructing approximate solutions to more difficult problems such as the traveling salesman problem (see Section 12.3).



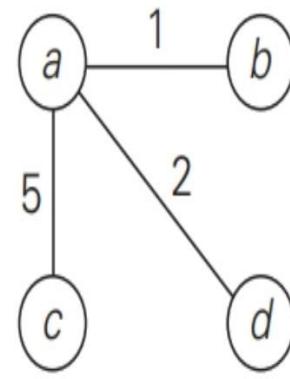
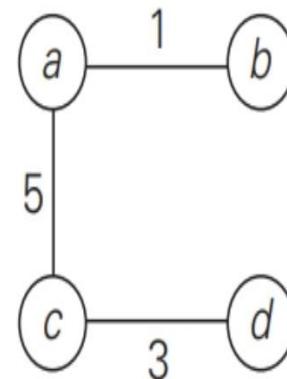
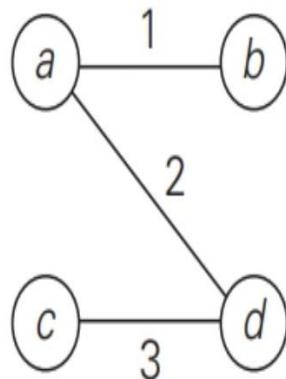
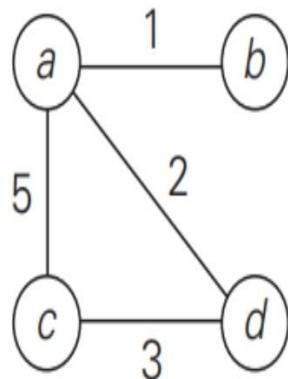
# Minimum Spanning Tree

We can represent the points given by vertices of a graph, possible connections by the graph's edges, and the connection costs by the edge weights. Then the question can be posed as the minimum spanning tree problem, defined formally as follows.

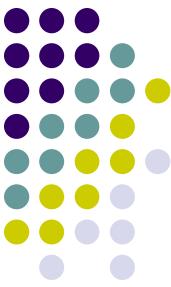
**DEFINITION** A *spanning tree* of an undirected connected graph is its connected acyclic subgraph (i.e., a tree) that contains all the vertices of the graph. If such a graph has weights assigned to its edges, a *minimum spanning tree* is its spanning tree of the smallest weight, where the *weight* of a tree is defined as the sum of the weights on all its edges. The *minimum spanning tree problem* is the problem of finding a minimum spanning tree for a given weighted connected graph.



# Minimum Spanning Tree

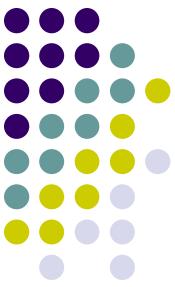


**FIGURE 9.2** Graph and its spanning trees, with  $T_1$  being the minimum spanning tree.



# Prim's Algorithm

Prim's algorithm constructs a minimum spanning tree through a sequence of expanding subtrees. The initial subtree in such a sequence consists of a single vertex selected arbitrarily from the set  $V$  of the graph's vertices. On each iteration, the algorithm expands the current tree in the greedy manner by simply attaching to it the nearest vertex not in that tree. (By the nearest vertex, we mean a vertex not in the tree connected to a vertex in the tree by an edge of the smallest weight. Ties can be broken arbitrarily.) The algorithm stops after all the graph's vertices have been included in the tree being constructed. Since the algorithm expands a tree by exactly one vertex on each of its iterations, the total number of such iterations is  $n - 1$ , where  $n$  is the number of vertices in the graph. The tree generated by the algorithm is obtained as the set of edges used for the tree expansions.



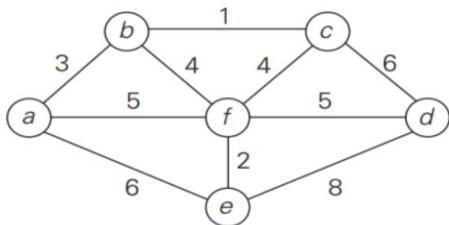
# Prim's Algorithm

**ALGORITHM** *Prim( $G$ )*

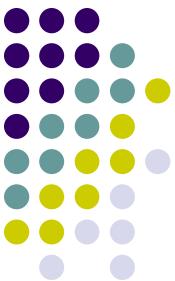
```
//Prim's algorithm for constructing a minimum spanning tree
//Input: A weighted connected graph  $G = \langle V, E \rangle$ 
//Output:  $E_T$ , the set of edges composing a minimum spanning tree of  $G$ 
 $V_T \leftarrow \{v_0\}$  //the set of tree vertices can be initialized with any vertex
 $E_T \leftarrow \emptyset$ 
for  $i \leftarrow 1$  to  $|V| - 1$  do
    find a minimum-weight edge  $e^* = (v^*, u^*)$  among all the edges  $(v, u)$ 
    such that  $v$  is in  $V_T$  and  $u$  is in  $V - V_T$ 
     $V_T \leftarrow V_T \cup \{u^*\}$ 
     $E_T \leftarrow E_T \cup \{e^*\}$ 
return  $E_T$ 
```



# Prim's Algorithm



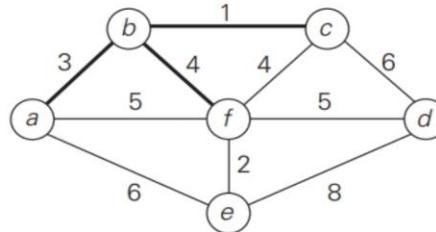
Tree vertices	Remaining vertices	Illustration
a(–, –)	<b>b(a, 3)</b> c(–, ∞) d(–, ∞) e(a, 6) f(a, 5)	
b(a, 3)	<b>c(b, 1)</b> d(–, ∞) e(a, 6) f(b, 4)	



# Prim's Algorithm

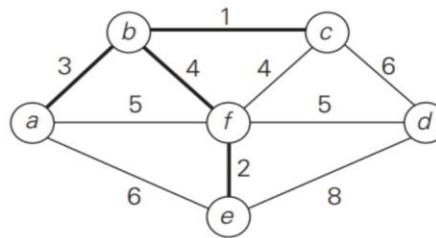
c(b, 1)

d(c, 6) e(a, 6) **f(b, 4)**



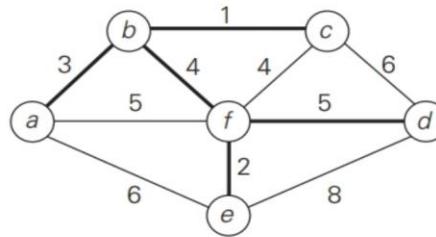
f(b, 4)

d(f, 5) **e(f, 2)**

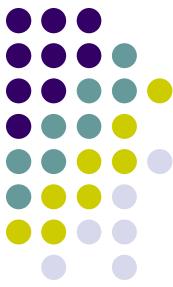


e(f, 2)

**d(f, 5)**



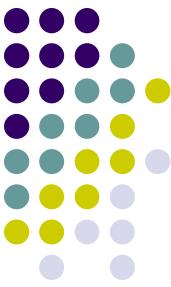
d(f, 5)



# Kruskal's Algorithm

he was a second-year graduate student [Kru56]. Kruskal's algorithm looks at a minimum spanning tree of a weighted connected graph  $G = \langle V, E \rangle$  as an acyclic subgraph with  $|V| - 1$  edges for which the sum of the edge weights is the smallest. (It is not difficult to prove that such a subgraph must be a tree.) Consequently, the algorithm constructs a minimum spanning tree as an expanding sequence of subgraphs that are always acyclic but are not necessarily connected on the intermediate stages of the algorithm.

The algorithm begins by sorting the graph's edges in nondecreasing order of their weights. Then, starting with the empty subgraph, it scans this sorted list, adding the next edge on the list to the current subgraph if such an inclusion does not create a cycle and simply skipping the edge otherwise.



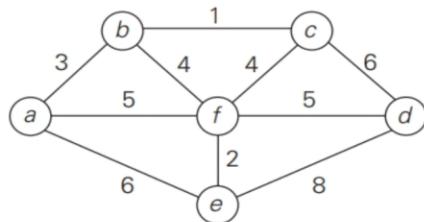
# Kruskal's Algorithm

**ALGORITHM** *Kruskal(G)*

```
//Kruskal's algorithm for constructing a minimum spanning tree
//Input: A weighted connected graph  $G = \langle V, E \rangle$ 
//Output:  $E_T$ , the set of edges composing a minimum spanning tree of  $G$ 
sort  $E$  in nondecreasing order of the edge weights  $w(e_{i_1}) \leq \dots \leq w(e_{i_{|E|}})$ 
 $E_T \leftarrow \emptyset$ ;  $e\text{counter} \leftarrow 0$       //initialize the set of tree edges and its size
 $k \leftarrow 0$                                 //initialize the number of processed edges
while  $e\text{counter} < |V| - 1$  do
     $k \leftarrow k + 1$ 
    if  $E_T \cup \{e_{i_k}\}$  is acyclic
         $E_T \leftarrow E_T \cup \{e_{i_k}\}$ ;  $e\text{counter} \leftarrow e\text{counter} + 1$ 
return  $E_T$ 
```



# Kruskal's Algorithm



---

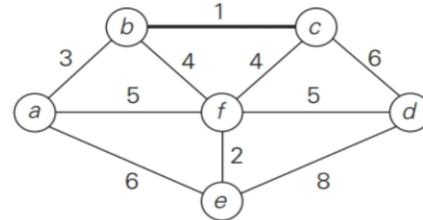
**Tree edges**

**Sorted list of edges**

**Illustration**

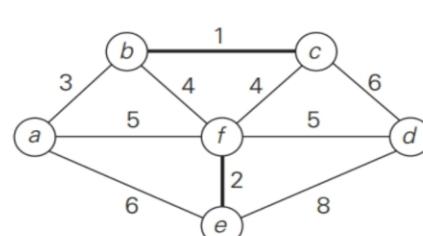
---

**bc**  
1    ef    ab    bf    cf    af    df    ae    cd    de



bc  
1

bc    **ef**    ab    bf    cf    af    df    ae    cd    de

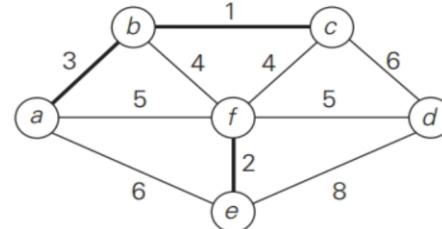




# Kruskal's Algorithm

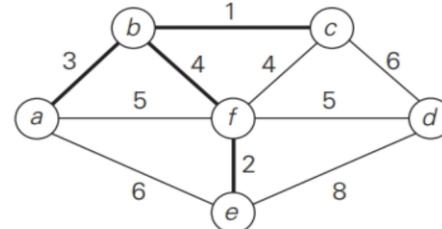
ef  
2

bc ef ab bf cf af df ae cd de  
1 2 3 4 4 4 5 5 6 6 8



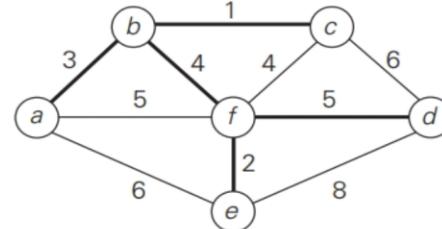
ab  
3

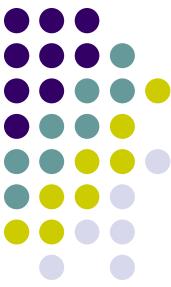
bc ef ab bf cf af df ae cd de  
1 2 3 4 4 4 5 5 6 6 8



bf  
4

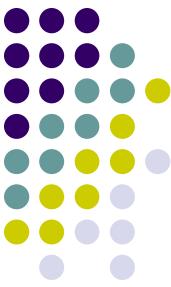
bc ef ab bf cf af df ae cd de  
1 2 3 4 4 4 5 5 6 6 8





# Dijkstra's Algorithm

In this section, we consider the ***single-source shortest-paths problem***: for a given vertex called the ***source*** in a weighted connected graph, find shortest paths to all its other vertices. It is important to stress that we are not interested here in a single shortest path that starts at the source and visits all the other vertices. This would have been a much more difficult problem (actually, a version of the traveling salesman problem introduced in Section 3.4 and discussed again later in the book). The single-source shortest-paths problem asks for a family of paths, each leading from the source to a different vertex in the graph, though some paths may, of course, have edges in common.



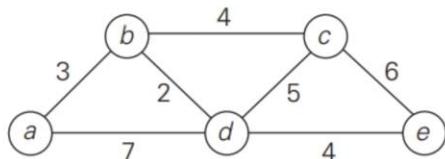
# Dijkstra's Algorithm

**ALGORITHM** *Dijkstra*( $G, s$ )

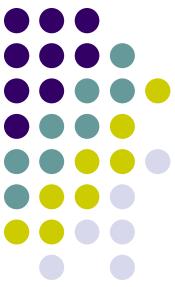
```
//Dijkstra's algorithm for single-source shortest paths
//Input: A weighted connected graph  $G = \langle V, E \rangle$  with nonnegative weights
//       and its vertex  $s$ 
//Output: The length  $d_v$  of a shortest path from  $s$  to  $v$ 
//       and its penultimate vertex  $p_v$  for every vertex  $v$  in  $V$ 
Initialize( $Q$ ) //initialize priority queue to empty
for every vertex  $v$  in  $V$ 
     $d_v \leftarrow \infty$ ;  $p_v \leftarrow \text{null}$ 
    Insert( $Q, v, d_v$ ) //initialize vertex priority in the priority queue
 $d_s \leftarrow 0$ ; Decrease( $Q, s, d_s$ ) //update priority of  $s$  with  $d_s$ 
 $V_T \leftarrow \emptyset$ 
for  $i \leftarrow 0$  to  $|V| - 1$  do
     $u^* \leftarrow \text{DeleteMin}(Q)$  //delete the minimum priority element
     $V_T \leftarrow V_T \cup \{u^*\}$ 
    for every vertex  $u$  in  $V - V_T$  that is adjacent to  $u^*$  do
        if  $d_{u^*} + w(u^*, u) < d_u$ 
             $d_u \leftarrow d_{u^*} + w(u^*, u)$ ;  $p_u \leftarrow u^*$ 
            Decrease( $Q, u, d_u$ )
```



# Dijkstra's Algorithm



Tree vertices	Remaining vertices	Illustration
a(−, 0)	<b>b(a, 3)</b> c(−, ∞) d(a, 7) e(−, ∞)	<pre>graph LR; a((a)) --- 7  d((d)); a --- 3  b((b)); a --- 2  d; b --- 4  c((c)); b --- 5  d; c --- 6  e((e)); d --- 4  e;</pre>
b(a, 3)	c(b, 3 + 4) <b>d(b, 3 + 2)</b> e(−, ∞)	<pre>graph LR; a((a)) --- 7  d((d)); a --- 3  b((b)); a --- 2  d; b --- 4  c((c)); b --- 5  d; c --- 6  e((e)); d --- 4  e;</pre>
d(b, 5)	<b>c(b, 7)</b> e(d, 5 + 4)	<pre>graph LR; a((a)) --- 7  d((d)); a --- 3  b((b)); a --- 2  d; b --- 4  c((c)); b --- 5  d; c --- 6  e((e)); d --- 4  e;</pre>

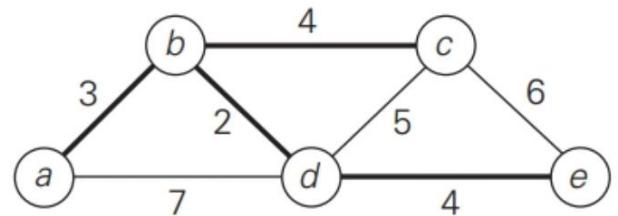


# Dijkstra's Algorithm

c(b, 7)

e(d, 9)

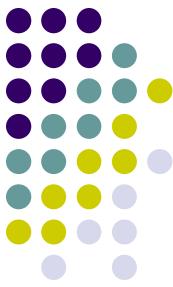
e(d, 9)





# Huffman Trees

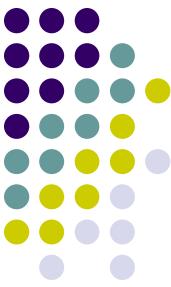
- Simple way to encode n symbols:
  - Assign a fixed length ( $m$ ) of bits to each symbols
  - $m \geq \log_2(n)$
- Improvement: assign more likely symbols shorter encodings:
  - Morse code:
    - E = .
    - O = - - -
  - AKA variable length encodings



# Huffman Trees

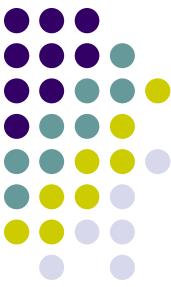
- Problem: how do we know when we have enough bits to decode a symbol?
- Solution: prefix-free codes.
  - No code word is prefix of another code word.

If we want to create a binary prefix code for some alphabet, it is natural to associate the alphabet's symbols with leaves of a binary tree in which all the left edges are labeled by 0 and all the right edges are labeled by 1. The codeword of a symbol can then be obtained by recording the labels on the simple path from the root to the symbol's leaf. Since there is no simple path to a leaf that continues to another leaf, no codeword can be a prefix of another codeword; hence, any such tree yields a prefix code.



# Huffman Trees

Among the many trees that can be constructed in this manner for a given alphabet with known frequencies of the symbol occurrences, how can we construct a tree that would assign shorter bit strings to high-frequency symbols and longer ones to low-frequency symbols? It can be done by the following greedy algorithm, invented by David Huffman while he was a graduate student at MIT [Huf52].

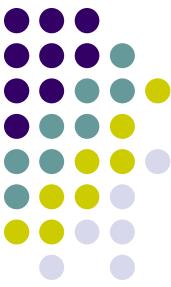


# Huffman Trees

## Huffman's algorithm

**Step 1** Initialize  $n$  one-node trees and label them with the symbols of the alphabet given. Record the frequency of each symbol in its tree's root to indicate the tree's ***weight***. (More generally, the weight of a tree will be equal to the sum of the frequencies in the tree's leaves.)

**Step 2** Repeat the following operation until a single tree is obtained. Find two trees with the smallest weight (ties can be broken arbitrarily, but see Problem 2 in this section's exercises). Make them the left and right subtree of a new tree and record the sum of their weights in the root of the new tree as its weight.



# Huffman Trees

## Huffman's algorithm

**Step 1** Initialize  $n$  one-node trees and label them with the symbols of the alphabet given. Record the frequency of each symbol in its tree's root to indicate the tree's ***weight***. (More generally, the weight of a tree will be equal to the sum of the frequencies in the tree's leaves.)

**Step 2** Repeat the following operation until a single tree is obtained. Find two trees with the smallest weight (ties can be broken arbitrarily, but see Problem 2 in this section's exercises). Make them the left and right subtree of a new tree and record the sum of their weights in the root of the new tree as its weight.



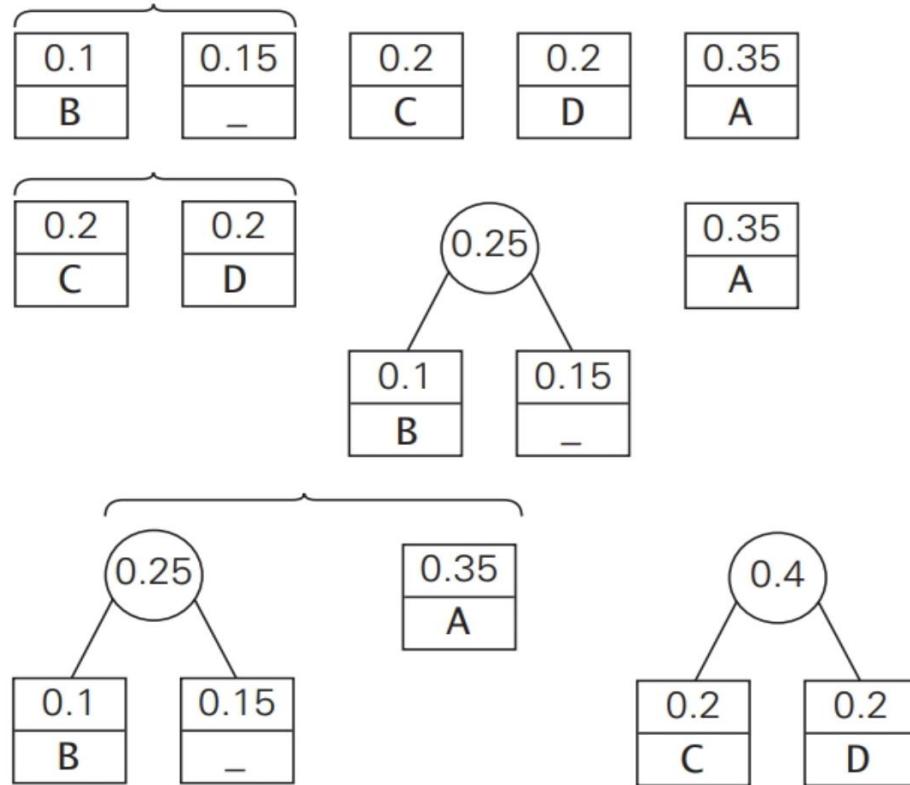
# Huffman Trees

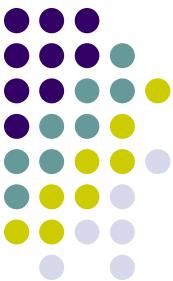
**EXAMPLE** Consider the five-symbol alphabet  $\{A, B, C, D, \_\}$  with the following occurrence frequencies in a text made up of these symbols:

symbol	A	B	C	D	_
frequency	0.35	0.1	0.2	0.2	0.15

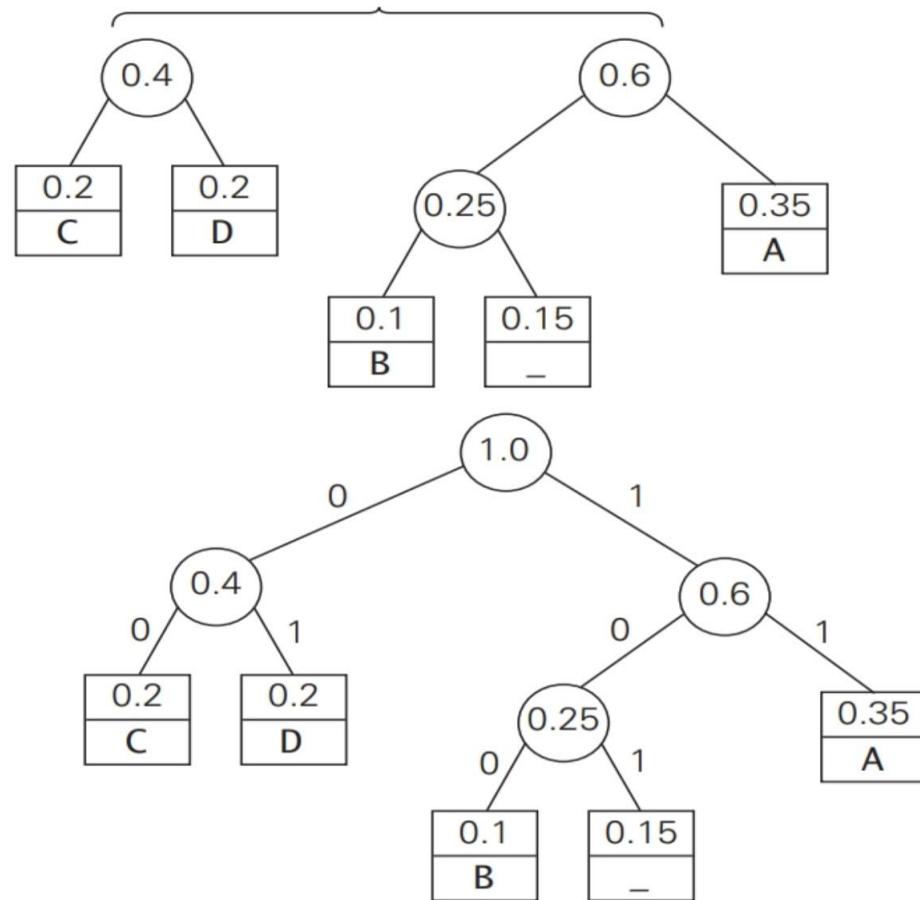


# Huffman Trees





# Huffman Trees

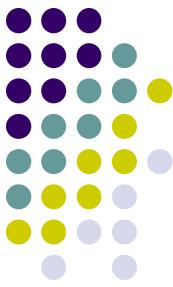




# Huffman Trees

The resulting codewords are as follows:

symbol	A	B	C	D	-
frequency	0.35	0.1	0.2	0.2	0.15
codeword	11	100	00	01	101

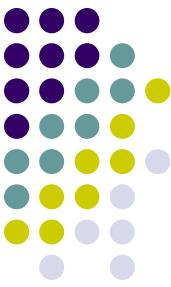


# Huffman Trees

With the occurrence frequencies given and the codeword lengths obtained, the average number of bits per symbol in this code is

$$2 \cdot 0.35 + 3 \cdot 0.1 + 2 \cdot 0.2 + 2 \cdot 0.2 + 3 \cdot 0.15 = 2.25.$$

Had we used a fixed-length encoding for the same alphabet, we would have to use at least 3 bits per each symbol. Thus, for this toy example, Huffman's code achieves the ***compression ratio***—a standard measure of a compression algorithm's effectiveness—of  $(3 - 2.25)/3 \cdot 100\% = 25\%$ . In other words, Huffman's encoding of the text will use 25% less memory than its fixed-length encoding. (Extensive experiments with Huffman codes have shown that the compression ratio for this scheme typically falls between 20% and 80%, depending on the characteristics of the text being compressed.) ■



# Activity Selection Problem

## An activity-selection problem

Our first example is the problem of scheduling several competing activities that require exclusive use of a common resource, with a goal of selecting a maximum-size set of mutually compatible activities. Suppose we have a set  $S = \{a_1, a_2, \dots, a_n\}$  of  $n$  proposed **activities** that wish to use a resource, such as a lecture hall, which can serve only one activity at a time. Each activity  $a_i$  has a **start time**  $s_i$  and a **finish time**  $f_i$ , where  $0 \leq s_i < f_i < \infty$ . If selected, activity  $a_i$  takes place during the half-open time interval  $[s_i, f_i)$ . Activities  $a_i$  and  $a_j$  are **compatible** if the intervals  $[s_i, f_i)$  and  $[s_j, f_j)$  do not overlap. That is,  $a_i$  and  $a_j$  are compatible if  $s_i \geq f_j$  or  $s_j \geq f_i$ . In the **activity-selection problem**, we wish to select a maximum-size subset of mutually compatible activities. We assume that the activities are sorted in monotonically increasing order of finish time:

$$f_1 \leq f_2 \leq f_3 \leq \dots \leq f_{n-1} \leq f_n . \quad (16.1)$$



# Activity Selection Problem

consider the following set  $S$  of activities:

$i$	1	2	3	4	5	6	7	8	9	10	11
$s_i$	1	3	0	5	3	5	6	8	8	2	12
$f_i$	4	5	6	7	9	9	10	11	12	14	16

For this example, the subset  $\{a_3, a_9, a_{11}\}$  consists of mutually compatible activities. It is not a maximum subset, however, since the subset  $\{a_1, a_4, a_8, a_{11}\}$  is larger. In fact,  $\{a_1, a_4, a_8, a_{11}\}$  is a largest subset of mutually compatible activities; another largest subset is  $\{a_2, a_4, a_9, a_{11}\}$ .



# Activity Selection Problem

## The optimal substructure of the activity-selection problem

We can easily verify that the activity-selection problem exhibits optimal substructure. Let us denote by  $S_{ij}$  the set of activities that start after activity  $a_i$  finishes and that finish before activity  $a_j$  starts. Suppose that we wish to find a maximum set of mutually compatible activities in  $S_{ij}$ , and suppose further that such a maximum set is  $A_{ij}$ , which includes some activity  $a_k$ . By including  $a_k$  in an optimal solution, we are left with two subproblems: finding mutually compatible activities in the set  $S_{ik}$  (activities that start after activity  $a_i$  finishes and that finish before activity  $a_k$  starts) and finding mutually compatible activities in the set  $S_{kj}$  (activities that start after activity  $a_k$  finishes and that finish before activity  $a_j$  starts). Let  $A_{ik} = A_{ij} \cap S_{ik}$  and  $A_{kj} = A_{ij} \cap S_{kj}$ , so that  $A_{ik}$  contains the activities in  $A_{ij}$  that finish before  $a_k$  starts and  $A_{kj}$  contains the activities in  $A_{ij}$  that start after  $a_k$  finishes. Thus, we have  $A_{ij} = A_{ik} \cup \{a_k\} \cup A_{kj}$ , and so the maximum-size set  $A_{ij}$  of mutually compatible activities in  $S_{ij}$  consists of  $|A_{ij}| = |A_{ik}| + |A_{kj}| + 1$  activities.

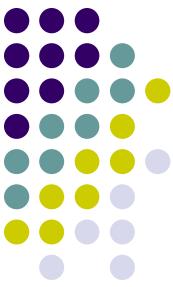
The usual cut-and-paste argument shows that the optimal solution  $A_{ij}$  must also include optimal solutions to the two subproblems for  $S_{ik}$  and  $S_{kj}$ . If we could find a set  $A'_{kj}$  of mutually compatible activities in  $S_{kj}$  where  $|A'_{kj}| > |A_{kj}|$ , then we could use  $A'_{kj}$ , rather than  $A_{kj}$ , in a solution to the subproblem for  $S_{ij}$ . We would have constructed a set of  $|A_{ik}| + |A'_{kj}| + 1 > |A_{ik}| + |A_{kj}| + 1 = |A_{ij}|$  mutually compatible activities, which contradicts the assumption that  $A_{ij}$  is an optimal solution. A symmetric argument applies to the activities in  $S_{ik}$ .



# Activity Selection Problem

This way of characterizing optimal substructure suggests that we might solve the activity-selection problem by dynamic programming. If we denote the size of an optimal solution for the set  $S_{ij}$  by  $c[i, j]$ , then we would have the recurrence

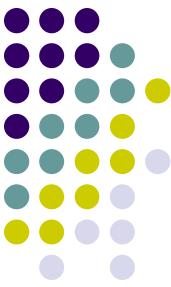
$$c[i, j] = c[i, k] + c[k, j] + 1.$$



# Activity Selection Problem

Of course, if we did not know that an optimal solution for the set  $S_{ij}$  includes activity  $a_k$ , we would have to examine all activities in  $S_{ij}$  to find which one to choose, so that

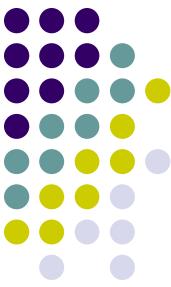
$$c[i, j] = \begin{cases} 0 & \text{if } S_{ij} = \emptyset , \\ \max_{a_k \in S_{ij}} \{c[i, k] + c[k, j] + 1\} & \text{if } S_{ij} \neq \emptyset . \end{cases} \quad (16.2)$$



# Activity Selection Problem

## Making the greedy choice

What if we could choose an activity to add to our optimal solution without having to first solve all the subproblems? That could save us from having to consider all the choices inherent in recurrence (16.2). In fact, for the activity-selection problem, we need consider only one choice: the greedy choice.



# Activity Selection Problem

What do we mean by the greedy choice for the activity-selection problem? Intuition suggests that we should choose an activity that leaves the resource available for as many other activities as possible. Now, of the activities we end up choosing, one of them must be the first one to finish. Our intuition tells us, therefore, to choose the activity in  $S$  with the earliest finish time, since that would leave the resource available for as many of the activities that follow it as possible. (If more than one activity in  $S$  has the earliest finish time, then we can choose any such activity.) In other words, since the activities are sorted in monotonically increasing order by finish time, the greedy choice is activity  $a_1$ . Choosing the first activity to finish is not the only way to think of making a greedy choice for this problem;



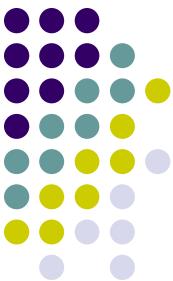
# Activity Selection Problem

If we make the greedy choice, we have only one remaining subproblem to solve: finding activities that start after  $a_1$  finishes. Why don't we have to consider activities that finish before  $a_1$  starts? We have that  $s_1 < f_1$ , and  $f_1$  is the earliest finish time of any activity, and therefore no activity can have a finish time less than or equal to  $s_1$ . Thus, all activities that are compatible with activity  $a_1$  must start after  $a_1$  finishes.



# Activity Selection Problem

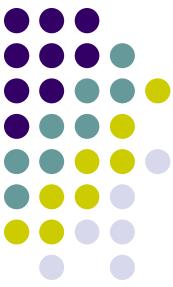
Furthermore, we have already established that the activity-selection problem exhibits optimal substructure. Let  $S_k = \{a_i \in S : s_i \geq f_k\}$  be the set of activities that start after activity  $a_k$  finishes. If we make the greedy choice of activity  $a_1$ , then  $S_1$  remains as the only subproblem to solve.<sup>1</sup> Optimal substructure tells us that if  $a_1$  is in the optimal solution, then an optimal solution to the original problem consists of activity  $a_1$  and all the activities in an optimal solution to the subproblem  $S_1$ .



# Activity Selection Problem

RECURSIVE-ACTIVITY-SELECTOR( $s, f, k, n$ )

- 1  $m = k + 1$
- 2 **while**  $m \leq n$  and  $s[m] < f[k]$  // find the first activity in  $S_k$  to finish
- 3      $m = m + 1$
- 4 **if**  $m \leq n$
- 5         **return**  $\{a_m\} \cup$  RECURSIVE-ACTIVITY-SELECTOR( $s, f, m, n$ )
- 6 **else return**  $\emptyset$



# Activity Selection Problem

**GREEDY-ACTIVITY-SELECTOR( $s, f$ )**

```
1   $n = s.length$ 
2   $A = \{a_1\}$ 
3   $k = 1$ 
4  for  $m = 2$  to  $n$ 
5      if  $s[m] \geq f[k]$ 
6           $A = A \cup \{a_m\}$ 
7           $k = m$ 
8  return  $A$ 
```



# Activity Selection Problem

